

bulldozer-price-prediction

September 12, 2023

1 Predicting sale price of Bulldozers using Machine Learning

1.1 1. Problem Definition

Predict the future sale price of a Bulldozer based on it's characteristics and past sales of similar bulldozers.

1.2 2. Data

The data is taken from the Kaggle competition - Blue Book for Bulldozers : <https://www.kaggle.com/competitions/bluebook-for-bulldozers/data>

There are 3 datasets for the problem - * Train.csv is the training set, which contains data through the end of 2011. * Valid.csv is the validation set, which contains data from January 1, 2012 - April 30, 2012. * Test.csv is the test set which contains data from May 1, 2012 - November 2012.

1.3 3. Evaluation

The evaluation metric for this competition is the RMSLE (root mean squared log error) between the actual and predicted auction prices.

Note - In a regression problem, the goal is to reduce the error as much as possible. i.e. for the given problem, reduce RMSLE.

1.4 4. Features

Kaggle provides a data dictionary for the features present in the datasets.

Data Dictionary google spreadsheet : https://docs.google.com/spreadsheets/d/1TdR-DKtUNUdwcuaWLoL_K8bA28_ms2z_7K5Hza0wng/edit?usp=sharing

Importing Modules

```
[60]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
from sklearn.ensemble import RandomForestRegressor

%matplotlib inline
```

1.4.1 Exploratory Data Analysis (EDA)

```
[2]: df = pd.read_csv("data/bluebook-for-bulldozers/TrainAndValid.csv", low_memory =  
    ↪False)
```

```
[3]: df.head()
```

```
[3]:   SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID  YearMade  \  
0    1139246    66000.0    999089    3157         121           3.0      2004  
1    1139248    57000.0    117657      77         121           3.0      1996  
2    1139249    10000.0    434808    7009         121           3.0      2001  
3    1139251    38500.0    1026470    332         121           3.0      2001  
4    1139253    11000.0    1057373   17311         121           3.0      2007
```

```
      MachineHoursCurrentMeter  UsageBand      saledate  ...  \  
0                68.0      Low  11/16/2006 0:00  ...  
1             4640.0      Low   3/26/2004 0:00  ...  
2             2838.0     High   2/26/2004 0:00  ...  
3             3486.0     High   5/19/2011 0:00  ...  
4              722.0   Medium   7/23/2009 0:00  ...
```

```
      Undercarriage_Pad_Width  Stick_Length  Thumb  Pattern_Changer  Grouser_Type  \  
0                NaN      NaN      NaN      NaN      NaN      NaN  
1                NaN      NaN      NaN      NaN      NaN      NaN  
2                NaN      NaN      NaN      NaN      NaN      NaN  
3                NaN      NaN      NaN      NaN      NaN      NaN  
4                NaN      NaN      NaN      NaN      NaN      NaN
```

```
      Backhoe_Mounting  Blade_Type  Travel_Controls  Differential_Type  \  
0                NaN      NaN      NaN      Standard  
1                NaN      NaN      NaN      Standard  
2                NaN      NaN      NaN      NaN  
3                NaN      NaN      NaN      NaN  
4                NaN      NaN      NaN      NaN
```

```
      Steering_Controls  
0      Conventional  
1      Conventional  
2                NaN  
3                NaN  
4                NaN
```

```
[5 rows x 53 columns]
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 412698 entries, 0 to 412697
```

Data columns (total 53 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	saledate	412698 non-null	object
10	fiModelDesc	412698 non-null	object
11	fiBaseModel	412698 non-null	object
12	fiSecondaryDesc	271971 non-null	object
13	fiModelSeries	58667 non-null	object
14	fiModelDescriptor	74816 non-null	object
15	ProductSize	196093 non-null	object
16	fiProductClassDesc	412698 non-null	object
17	state	412698 non-null	object
18	ProductGroup	412698 non-null	object
19	ProductGroupDesc	412698 non-null	object
20	Drive_System	107087 non-null	object
21	Enclosure	412364 non-null	object
22	Forks	197715 non-null	object
23	Pad_Type	81096 non-null	object
24	Ride_Control	152728 non-null	object
25	Stick	81096 non-null	object
26	Transmission	188007 non-null	object
27	Turbocharged	81096 non-null	object
28	Blade_Extension	25983 non-null	object
29	Blade_Width	25983 non-null	object
30	Enclosure_Type	25983 non-null	object
31	Engine_Horsepower	25983 non-null	object
32	Hydraulics	330133 non-null	object
33	Pushblock	25983 non-null	object
34	Ripper	106945 non-null	object
35	Scarifier	25994 non-null	object
36	Tip_Control	25983 non-null	object
37	Tire_Size	97638 non-null	object
38	Coupler	220679 non-null	object
39	Coupler_System	44974 non-null	object
40	Grouser_Tracks	44875 non-null	object
41	Hydraulics_Flow	44875 non-null	object
42	Track_Type	102193 non-null	object
43	Undercarriage_Pad_Width	102916 non-null	object
44	Stick_Length	102261 non-null	object

```

45 Thumb 102332 non-null object
46 Pattern_Changer 102261 non-null object
47 Grouser_Type 102193 non-null object
48 Backhoe_Mounting 80712 non-null object
49 Blade_Type 81875 non-null object
50 Travel_Controls 81877 non-null object
51 Differential_Type 71564 non-null object
52 Steering_Controls 71522 non-null object
dtypes: float64(3), int64(5), object(45)
memory usage: 166.9+ MB

```

```
[5]: df.isna().sum()
```

```

[5]: SalesID 0
SalePrice 0
MachineID 0
ModelID 0
datasource 0
auctioneerID 20136
YearMade 0
MachineHoursCurrentMeter 265194
UsageBand 339028
saledate 0
fiModelDesc 0
fiBaseModel 0
fiSecondaryDesc 140727
fiModelSeries 354031
fiModelDescriptor 337882
ProductSize 216605
fiProductClassDesc 0
state 0
ProductGroup 0
ProductGroupDesc 0
Drive_System 305611
Enclosure 334
Forks 214983
Pad_Type 331602
Ride_Control 259970
Stick 331602
Transmission 224691
Turbocharged 331602
Blade_Extension 386715
Blade_Width 386715
Enclosure_Type 386715
Engine_Horsepower 386715
Hydraulics 82565
Pushblock 386715

```

Ripper	305753
Scarifier	386704
Tip_Control	386715
Tire_Size	315060
Coupler	192019
Coupler_System	367724
Grouser_Tracks	367823
Hydraulics_Flow	367823
Track_Type	310505
Undercarriage_Pad_Width	309782
Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176

dtype: int64

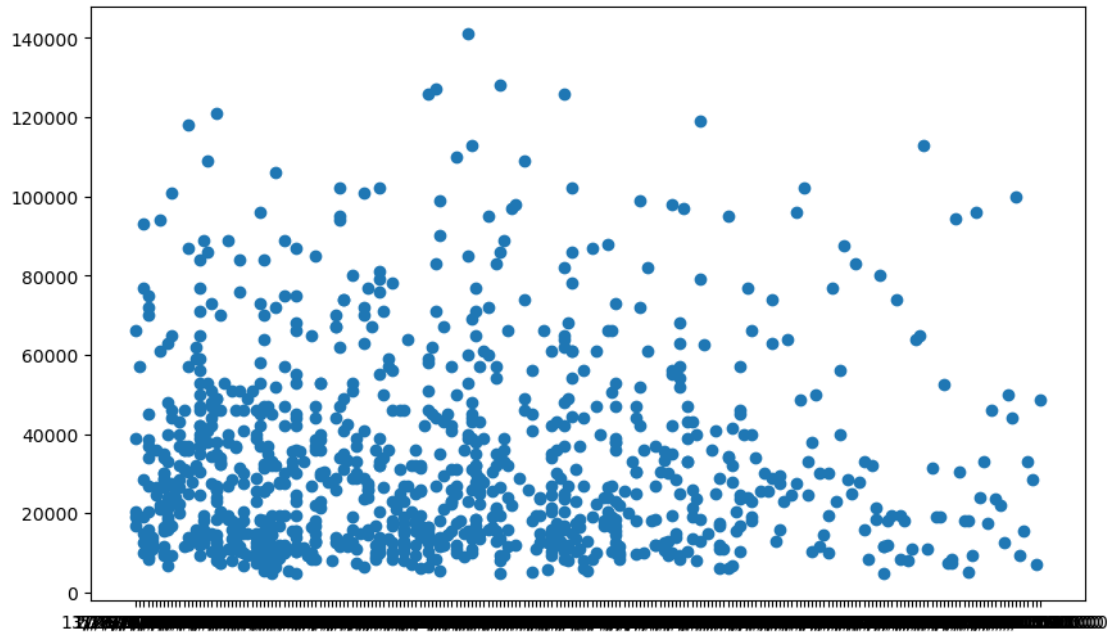
```
[6]: df["saledate"][:10]
```

```
[6]: 0    11/16/2006 0:00
     1     3/26/2004 0:00
     2     2/26/2004 0:00
     3     5/19/2011 0:00
     4     7/23/2009 0:00
     5    12/18/2008 0:00
     6     8/26/2004 0:00
     7    11/17/2005 0:00
     8     8/27/2009 0:00
     9     8/9/2007 0:00
     Name: saledate, dtype: object
```

```
[7]: df["saledate"].dtype
```

```
[7]: dtype('O')
```

```
[8]: fig, ax = plt.subplots(figsize = (10,6))
     ax = plt.scatter(df["saledate"][:1000], df["SalePrice"][:1000]);
```



The saledate column contains datetime data and must be parsed as datetime during importation

```
[9]: # Reimporting with parsed datetime
df = pd.read_csv("data/bluebook-for-bulldozers/TrainAndValid.csv",
                 low_memory=False, parse_dates=["saledate"])
```

```
[10]: df.head()
```

```
[10]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	\
0	1139246	66000.0	999089	3157	121	3.0	2004	
1	1139248	57000.0	117657	77	121	3.0	1996	
2	1139249	10000.0	434808	7009	121	3.0	2001	
3	1139251	38500.0	1026470	332	121	3.0	2001	
4	1139253	11000.0	1057373	17311	121	3.0	2007	

	MachineHoursCurrentMeter	UsageBand	saledate	...	Undercarriage_Pad_Width	\
0	68.0	Low	2006-11-16	...	NaN	
1	4640.0	Low	2004-03-26	...	NaN	
2	2838.0	High	2004-02-26	...	NaN	
3	3486.0	High	2011-05-19	...	NaN	
4	722.0	Medium	2009-07-23	...	NaN	

	Stick_Length	Thumb	Pattern_Changer	Grouser_Type	Backhoe_Mounting	Blade_Type	\
0	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	NaN	

3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN

	Travel_Controls	Differential_Type	Steering_Controls
0	NaN	Standard	Conventional
1	NaN	Standard	Conventional
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 53 columns]

```
[11]: df["saledate"][:10]
```

```
[11]: 0    2006-11-16
      1    2004-03-26
      2    2004-02-26
      3    2011-05-19
      4    2009-07-23
      5    2008-12-18
      6    2004-08-26
      7    2005-11-17
      8    2009-08-27
      9    2007-08-09
      Name: saledate, dtype: datetime64[ns]
```

```
[12]: df["saledate"].dtype
```

```
[12]: dtype('<M8[ns]')
```

```
[13]: fig, ax = plt.subplots(figsize = (10,6))
      ax = plt.scatter(df["saledate"][:1000], df["SalePrice"][:1000])
```



Since Jupyter Notebook truncates the data if the number of features is large let's view all the features in `head()` using it's transpose.

```
[14]: df.head().T
```

```
[14]:
```

SalesID	1139246
SalePrice	66000.0
MachineID	999089
ModelID	3157
datasource	121
auctioneerID	3.0
YearMade	2004
MachineHoursCurrentMeter	68.0
UsageBand	Low
saledate	2006-11-16 00:00:00
fiModelDesc	521D
fiBaseModel	521
fiSecondaryDesc	D
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - 110.0 to 120.0 Horsepower
state	Alabama
ProductGroup	WL
ProductGroupDesc	Wheel Loader

Drive_System	NaN
Enclosure	EROPS w AC
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional

	1 \
SalesID	1139248
SalePrice	57000.0
MachineID	117657
ModelID	77
datasource	121
auctioneerID	3.0
YearMade	1996
MachineHoursCurrentMeter	4640.0
UsageBand	Low
saledate	2004-03-26 00:00:00
fiModelDesc	950FII
fiBaseModel	950

fiSecondaryDesc	F
fiModelSeries	II
fiModelDescriptor	NaN
ProductSize	Medium
fiProductClassDesc	Wheel Loader - 150.0 to 175.0 Horsepower
state	North Carolina
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	EROPS w AC
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	23.5
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
2 \	
SalesID	1139249
SalePrice	10000.0
MachineID	434808
ModelID	7009

datasource	121
auctioneerID	3.0
YearMade	2001
MachineHoursCurrentMeter	2838.0
UsageBand	High
saledate	2004-02-26 00:00:00
fiModelDesc	226
fiBaseModel	226
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...
state	New York
ProductGroup	SSL
ProductGroupDesc	Skid Steer Loaders
Drive_System	NaN
Enclosure	OROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	Auxiliary
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	None or Unspecified
Coupler_System	None or Unspecified
Grouser_Tracks	None or Unspecified
Hydraulics_Flow	Standard
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN

Differential_Type		NaN
Steering_Controls		NaN
		3 \
SalesID		1139251
SalePrice		38500.0
MachineID		1026470
ModelID		332
datasource		121
auctioneerID		3.0
YearMade		2001
MachineHoursCurrentMeter		3486.0
UsageBand		High
saledate	2011-05-19 00:00:00	
fiModelDesc		PC120-6E
fiBaseModel		PC120
fiSecondaryDesc		NaN
fiModelSeries		-6E
fiModelDescriptor		NaN
ProductSize		Small
fiProductClassDesc	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...	
state		Texas
ProductGroup		TEX
ProductGroupDesc	Track Excavators	
Drive_System		NaN
Enclosure		EROPS w AC
Forks		NaN
Pad_Type		NaN
Ride_Control		NaN
Stick		NaN
Transmission		NaN
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics		2 Valve
Pushblock		NaN
Ripper		NaN
Scarifier		NaN
Tip_Control		NaN
Tire_Size		NaN
Coupler	None or Unspecified	
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN

Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	NaN
Steering_Controls	NaN
	4
SalesID	1139253
SalePrice	11000.0
MachineID	1057373
ModelID	17311
datasource	121
auctioneerID	3.0
YearMade	2007
MachineHoursCurrentMeter	722.0
UsageBand	Medium
saledate	2009-07-23 00:00:00
fiModelDesc	S175
fiBaseModel	S175
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Skid Steer Loader - 1601.0 to 1751.0 Lb Operat...
state	New York
ProductGroup	SSL
ProductGroupDesc	Skid Steer Loaders
Drive_System	NaN
Enclosure	EROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	Auxiliary
Pushblock	NaN
Ripper	NaN

Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	None or Unspecified
Coupler_System	None or Unspecified
Grouser_Tracks	None or Unspecified
Hydraulics_Flow	Standard
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	NaN
Steering_Controls	NaN

Since we are dealing with a timeseries data, it's better to sort the data according to the saledate

```
[15]: df["saledate"].head(10)
```

```
[15]: 0    2006-11-16
      1    2004-03-26
      2    2004-02-26
      3    2011-05-19
      4    2009-07-23
      5    2008-12-18
      6    2004-08-26
      7    2005-11-17
      8    2009-08-27
      9    2007-08-09
      Name: saledate, dtype: datetime64[ns]
```

1.4.2 Sorting the Dataframe by Saledate

```
[16]: df.sort_values(by = ["saledate"], inplace = True)
```

```
[17]: df["saledate"].head(10)
```

```
[17]: 205615    1989-01-17
      274835    1989-01-31
      141296    1989-01-31
      212552    1989-01-31
      62755    1989-01-31
      54653    1989-01-31
      81383    1989-01-31
```

```

204924    1989-01-31
135376    1989-01-31
113390    1989-01-31
Name: saledate, dtype: datetime64[ns]

```

```
[18]: df.head()
```

```

[18]:      SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID  \
205615  1646770    9500.0    1126363    8434          132          18.0
274835  1821514   14000.0    1194089    10150          132          99.0
141296  1505138   50000.0    1473654     4139          132          99.0
212552  1671174   16000.0    1327630     8591          132          99.0
62755   1329056   22000.0    1336053     4089          132          99.0

```

```

      YearMade  MachineHoursCurrentMeter  UsageBand  saledate  ...  \
205615      1974                      NaN         NaN  1989-01-17  ...
274835      1980                      NaN         NaN  1989-01-31  ...
141296      1978                      NaN         NaN  1989-01-31  ...
212552      1980                      NaN         NaN  1989-01-31  ...
62755      1984                      NaN         NaN  1989-01-31  ...

```

```

      Undercarriage_Pad_Width  Stick_Length  Thumb  Pattern_Changer  \
205615                      NaN          NaN    NaN              NaN
274835                      NaN          NaN    NaN              NaN
141296                      NaN          NaN    NaN              NaN
212552                      NaN          NaN    NaN              NaN
62755                       NaN          NaN    NaN              NaN

```

```

      Grouser_Type  Backhoe_Mounting  Blade_Type  Travel_Controls  \
205615          NaN  None or Unspecified  Straight  None or Unspecified
274835          NaN                NaN        NaN              NaN
141296          NaN  None or Unspecified  Straight  None or Unspecified
212552          NaN                NaN        NaN              NaN
62755          NaN  None or Unspecified        PAT              Lever

```

```

      Differential_Type  Steering_Controls
205615                NaN                NaN
274835          Standard          Conventional
141296                NaN                NaN
212552          Standard          Conventional
62755                NaN                NaN

```

```
[5 rows x 53 columns]
```

Now let's make a copy of the data for further modifications.

Making a copy of the DataFrame

```
[19]: df_tmp = df.copy()
```

```
[20]: df_tmp.head()
```

```
[20]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	\
205615	1646770	9500.0	1126363	8434	132	18.0	
274835	1821514	14000.0	1194089	10150	132	99.0	
141296	1505138	50000.0	1473654	4139	132	99.0	
212552	1671174	16000.0	1327630	8591	132	99.0	
62755	1329056	22000.0	1336053	4089	132	99.0	

	YearMade	MachineHoursCurrentMeter	UsageBand	saledate	...	\
205615	1974		NaN	NaN 1989-01-17	...	
274835	1980		NaN	NaN 1989-01-31	...	
141296	1978		NaN	NaN 1989-01-31	...	
212552	1980		NaN	NaN 1989-01-31	...	
62755	1984		NaN	NaN 1989-01-31	...	

	Undercarriage_Pad_Width	Stick_Length	Thumb	Pattern_Changer	\
205615		NaN	NaN	NaN	
274835		NaN	NaN	NaN	
141296		NaN	NaN	NaN	
212552		NaN	NaN	NaN	
62755		NaN	NaN	NaN	

	Grouser_Type	Backhoe_Mounting	Blade_Type	Travel_Controls	\
205615	NaN	None or Unspecified	Straight	None or Unspecified	
274835	NaN	NaN	NaN	NaN	
141296	NaN	None or Unspecified	Straight	None or Unspecified	
212552	NaN	NaN	NaN	NaN	
62755	NaN	None or Unspecified	PAT	Lever	

	Differential_Type	Steering_Controls
205615	NaN	NaN
274835	Standard	Conventional
141296	NaN	NaN
212552	Standard	Conventional
62755	NaN	NaN

```
[5 rows x 53 columns]
```

1.4.3 Feature engineering (Enriching data using the saledate feature)

```
[21]: df_tmp["saleYear"] = df_tmp["saledate"].dt.year
df_tmp["saleMonth"] = df_tmp["saledate"].dt.month
df_tmp["saleDay"] = df_tmp["saledate"].dt.day
df_tmp["saleDayOfWeek"] = df_tmp["saledate"].dt.dayofweek
df_tmp["saleDayOfYear"] = df_tmp["saledate"].dt.dayofyear
df_tmp.head().T
```


[21]:		205615 \
SalesID		1646770
SalePrice		9500.0
MachineID		1126363
ModelID		8434
datasource		132
auctioneerID		18.0
YearMade		1974
MachineHoursCurrentMeter		NaN
UsageBand		NaN
saledate		1989-01-17 00:00:00
fiModelDesc		TD20
fiBaseModel		TD20
fiSecondaryDesc		NaN
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		Medium
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	
state		Texas
ProductGroup		TTT
ProductGroupDesc	Track Type Tractors	
Drive_System		NaN
Enclosure		OROPS
Forks		NaN
Pad_Type		NaN
Ride_Control		NaN
Stick		NaN
Transmission	Direct Drive	
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics	2 Valve	
Pushblock		NaN
Ripper	None or Unspecified	
Scarifier		NaN
Tip_Control		NaN
Tire_Size		NaN
Coupler		NaN
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN

Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting	None or Unspecified	
Blade_Type	Straight	
Travel_Controls	None or Unspecified	
Differential_Type		NaN
Steering_Controls		NaN
saleYear		1989
saleMonth		1
saleDay		17
saleDayOfWeek		1
saleDayOfYear		17
	274835 \	
SalesID	1821514	
SalePrice	14000.0	
MachineID	1194089	
ModelID	10150	
datasource	132	
auctioneerID	99.0	
YearMade	1980	
MachineHoursCurrentMeter	NaN	
UsageBand	NaN	
saledate	1989-01-31 00:00:00	
fiModelDesc	A66	
fiBaseModel	A66	
fiSecondaryDesc	NaN	
fiModelSeries	NaN	
fiModelDescriptor	NaN	
ProductSize	NaN	
fiProductClassDesc	Wheel Loader - 120.0 to 135.0 Horsepower	
state	Florida	
ProductGroup	WL	
ProductGroupDesc	Wheel Loader	
Drive_System	NaN	
Enclosure	OROPS	
Forks	None or Unspecified	
Pad_Type	NaN	
Ride_Control	None or Unspecified	
Stick	NaN	
Transmission	NaN	
Turbocharged	NaN	
Blade_Extension	NaN	
Blade_Width	NaN	
Enclosure_Type	NaN	
Engine_Horsepower	NaN	
Hydraulics	2 Valve	

Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31
	141296 \
SalesID	1505138
SalePrice	50000.0
MachineID	1473654
ModelID	4139
datasource	132
auctioneerID	99.0
YearMade	1978
MachineHoursCurrentMeter	NaN
UsageBand	NaN
saledate	1989-01-31 00:00:00
fiModelDesc	D7G
fiBaseModel	D7
fiSecondaryDesc	G
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	Large
fiProductClassDesc	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...
state	Florida
ProductGroup	TTT
ProductGroupDesc	Track Type Tractors

Drive_System	NaN
Enclosure	OROPS
Forks	NaN
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	Standard
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

	212552 \
SalesID	1671174
SalePrice	16000.0
MachineID	1327630
ModelID	8591
datasource	132
auctioneerID	99.0
YearMade	1980

MachineHoursCurrentMeter		NaN
UsageBand		NaN
saledate	1989-01-31 00:00:00	
fiModelDesc		A62
fiBaseModel		A62
fiSecondaryDesc		NaN
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Wheel Loader - Unidentified	
state	Florida	
ProductGroup	WL	
ProductGroupDesc	Wheel Loader	
Drive_System		NaN
Enclosure		EROPS
Forks	None or Unspecified	
Pad_Type		NaN
Ride_Control	None or Unspecified	
Stick		NaN
Transmission		NaN
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics	2 Valve	
Pushblock		NaN
Ripper		NaN
Scarifier		NaN
Tip_Control		NaN
Tire_Size	None or Unspecified	
Coupler	None or Unspecified	
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		NaN
Blade_Type		NaN
Travel_Controls		NaN
Differential_Type	Standard	
Steering_Controls	Conventional	
saleYear		1989

saleMonth	1	
saleDay	31	
saleDayOfWeek	1	
saleDayOfYear	31	
		62755
SalesID		1329056
SalePrice		22000.0
MachineID		1336053
ModelID		4089
datasource		132
auctioneerID		99.0
YearMade		1984
MachineHoursCurrentMeter		NaN
UsageBand		NaN
saledate	1989-01-31 00:00:00	
fiModelDesc		D3B
fiBaseModel		D3
fiSecondaryDesc		B
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...	
state		Florida
ProductGroup		TTT
ProductGroupDesc	Track Type Tractors	
Drive_System		NaN
Enclosure		OROPS
Forks		NaN
Pad_Type		NaN
Ride_Control		NaN
Stick		NaN
Transmission		Standard
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics		2 Valve
Pushblock		NaN
Ripper	None or Unspecified	
Scarifier		NaN
Tip_Control		NaN
Tire_Size		NaN
Coupler		NaN
Coupler_System		NaN
Grouser_Tracks		NaN

Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	PAT
Travel_Controls	Lever
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

Now we don't need saledate column so let's drop it. If required, it is present in the original copy of DataFrame df.

```
[22]: df_tmp.drop("saledate", axis = 1, inplace = True)
df_tmp.head().T
```

```
[22]:
```

	205615	\
SalesID	1646770	
SalePrice	9500.0	
MachineID	1126363	
ModelID	8434	
datasource	132	
auctioneerID	18.0	
YearMade	1974	
MachineHoursCurrentMeter	NaN	
UsageBand	NaN	
fiModelDesc	TD20	
fiBaseModel	TD20	
fiSecondaryDesc	NaN	
fiModelSeries	NaN	
fiModelDescriptor	NaN	
ProductSize	Medium	
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	
state	Texas	
ProductGroup	TTT	
ProductGroupDesc	Track Type Tractors	
Drive_System	NaN	
Enclosure	OROPS	
Forks	NaN	
Pad_Type	NaN	

Ride_Control	NaN
Stick	NaN
Transmission	Direct Drive
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	17
saleDayOfWeek	1
saleDayOfYear	17
	274835 \
SalesID	1821514
SalePrice	14000.0
MachineID	1194089
ModelID	10150
datasource	132
auctioneerID	99.0
YearMade	1980
MachineHoursCurrentMeter	NaN
UsageBand	NaN
fiModelDesc	A66
fiBaseModel	A66

fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - 120.0 to 135.0 Horsepower
state	Florida
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	OROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

	141296	\
SalesID	1505138	
SalePrice	50000.0	
MachineID	1473654	
ModelID	4139	
datasource	132	
auctioneerID	99.0	
YearMade	1978	
MachineHoursCurrentMeter	NaN	
UsageBand	NaN	
fiModelDesc	D7G	
fiBaseModel	D7	
fiSecondaryDesc	G	
fiModelSeries	NaN	
fiModelDescriptor	NaN	
ProductSize	Large	
fiProductClassDesc	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	
state	Florida	
ProductGroup	TTT	
ProductGroupDesc	Track Type Tractors	
Drive_System	NaN	
Enclosure	OROPS	
Forks	NaN	
Pad_Type	NaN	
Ride_Control	NaN	
Stick	NaN	
Transmission	Standard	
Turbocharged	NaN	
Blade_Extension	NaN	
Blade_Width	NaN	
Enclosure_Type	NaN	
Engine_Horsepower	NaN	
Hydraulics	2 Valve	
Pushblock	NaN	
Ripper	None or Unspecified	
Scarifier	NaN	
Tip_Control	NaN	
Tire_Size	NaN	
Coupler	NaN	
Coupler_System	NaN	
Grouser_Tracks	NaN	
Hydraulics_Flow	NaN	
Track_Type	NaN	
Undercarriage_Pad_Width	NaN	
Stick_Length	NaN	
Thumb	NaN	
Pattern_Changer	NaN	

Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

	212552 \
SalesID	1671174
SalePrice	16000.0
MachineID	1327630
ModelID	8591
datasource	132
auctioneerID	99.0
YearMade	1980
MachineHoursCurrentMeter	NaN
UsageBand	NaN
fiModelDesc	A62
fiBaseModel	A62
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - Unidentified
state	Florida
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	EROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN

Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31
SalesID	62755
SalePrice	1329056
MachineID	22000.0
ModelID	1336053
datasource	4089
auctioneerID	132
YearMade	99.0
MachineHoursCurrentMeter	1984
UsageBand	NaN
fiModelDesc	NaN
fiBaseModel	D3B
fiSecondaryDesc	D3
fiModelSeries	B
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...
state	Florida
ProductGroup	TTT
ProductGroupDesc	Track Type Tractors
Drive_System	NaN
Enclosure	OROPS
Forks	NaN

Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	Standard
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	PAT
Travel_Controls	Lever
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

Let's explore the sales by US States

```
[23]: df_tmp.value_counts("state")
```

```
[23]: state
Florida      67320
Texas        53110
California   29761
Washington   16222
Georgia      14633
Maryland     13322
Mississippi  13240
```

Ohio	12369
Illinois	11540
Colorado	11529
New Jersey	11156
North Carolina	10636
Tennessee	10298
Alabama	10292
Pennsylvania	10234
South Carolina	9951
Arizona	9364
New York	8639
Connecticut	8276
Minnesota	7885
Missouri	7178
Nevada	6932
Louisiana	6627
Kentucky	5351
Maine	5096
Indiana	4124
Arkansas	3933
New Mexico	3631
Utah	3046
Unspecified	2801
Wisconsin	2745
New Hampshire	2738
Virginia	2353
Idaho	2025
Oregon	1911
Michigan	1831
Wyoming	1672
Montana	1336
Iowa	1336
Oklahoma	1326
Nebraska	866
West Virginia	840
Kansas	667
Delaware	510
North Dakota	480
Alaska	430
Massachusetts	347
Vermont	300
South Dakota	244
Hawaii	118
Rhode Island	83
Puerto Rico	42
Washington DC	2
dtype: int64	

```
[24]: df_tmp.value_counts("ProductGroupDesc")
```

```
[24]: ProductGroupDesc
Track Excavators      104230
Track Type Tractors   82582
Backhoe Loaders       81401
Wheel Loader          73216
Skid Steer Loaders    45011
Motor Graders         26258
dtype: int64
```

Before we go ahead with modelling we need to convert `string` data into pandas `categories`.

1.4.4 Converting string data to pandas categories

```
[25]: pd.api.types.is_string_dtype(df["state"])
```

```
[25]: True
```

```
[26]: df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SalesID                              412698 non-null  int64
1   SalePrice                            412698 non-null  float64
2   MachineID                            412698 non-null  int64
3   ModelID                              412698 non-null  int64
4   datasource                           412698 non-null  int64
5   auctioneerID                         392562 non-null  float64
6   YearMade                             412698 non-null  int64
7   MachineHoursCurrentMeter             147504 non-null  float64
8   UsageBand                            73670 non-null   object
9   fiModelDesc                           412698 non-null  object
10  fiBaseModel                           412698 non-null  object
11  fiSecondaryDesc                       271971 non-null  object
12  fiModelSeries                         58667 non-null   object
13  fiModelDescriptor                     74816 non-null   object
14  ProductSize                           196093 non-null  object
15  fiProductClassDesc                   412698 non-null  object
16  state                                 412698 non-null  object
17  ProductGroup                         412698 non-null  object
18  ProductGroupDesc                     412698 non-null  object
19  Drive_System                         107087 non-null  object
20  Enclosure                             412364 non-null  object
21  Forks                                197715 non-null  object
```

22	Pad_Type	81096	non-null	object
23	Ride_Control	152728	non-null	object
24	Stick	81096	non-null	object
25	Transmission	188007	non-null	object
26	Turbocharged	81096	non-null	object
27	Blade_Extension	25983	non-null	object
28	Blade_Width	25983	non-null	object
29	Enclosure_Type	25983	non-null	object
30	Engine_Horsepower	25983	non-null	object
31	Hydraulics	330133	non-null	object
32	Pushblock	25983	non-null	object
33	Ripper	106945	non-null	object
34	Scarifier	25994	non-null	object
35	Tip_Control	25983	non-null	object
36	Tire_Size	97638	non-null	object
37	Coupler	220679	non-null	object
38	Coupler_System	44974	non-null	object
39	Grouser_Tracks	44875	non-null	object
40	Hydraulics_Flow	44875	non-null	object
41	Track_Type	102193	non-null	object
42	Undercarriage_Pad_Width	102916	non-null	object
43	Stick_Length	102261	non-null	object
44	Thumb	102332	non-null	object
45	Pattern_Changer	102261	non-null	object
46	Grouser_Type	102193	non-null	object
47	Backhoe_Mounting	80712	non-null	object
48	Blade_Type	81875	non-null	object
49	Travel_Controls	81877	non-null	object
50	Differential_Type	71564	non-null	object
51	Steering_Controls	71522	non-null	object
52	saleYear	412698	non-null	int64
53	saleMonth	412698	non-null	int64
54	saleDay	412698	non-null	int64
55	saleDayOfWeek	412698	non-null	int64
56	saleDayOfYear	412698	non-null	int64

dtypes: float64(3), int64(10), object(44)

memory usage: 182.6+ MB

Let's iterate through the dataframe and check which features have string type values

```
[27]: for col_name , col_content in df_tmp.items():
        if pd.api.types.is_string_dtype(col_content):
            print(col_name)
```

```
UsageBand
fiModelDesc
fiBaseModel
fiSecondaryDesc
fiModelSeries
```



```

fiModelDescriptor
ProductSize
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
Tire_Size
Coupler
Coupler_System
Grouser_Tracks
Hydraulics_Flow
Track_Type
Undercarriage_Pad_Width
Stick_Length
Thumb
Pattern_Changer
Grouser_Type
Backhoe_Mounting
Blade_Type
Travel_Controls
Differential_Type
Steering_Controls

```

Let's change the data type of these features to pandas categories

```

[28]: for col_name, col_content in df_tmp.items():
        if pd.api.types.is_string_dtype(col_content):
            df_tmp[col_name] = col_content.astype("category").cat.as_ordered()
df_tmp.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203

```

Data columns (total 57 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	category
9	fiModelDesc	412698 non-null	category
10	fiBaseModel	412698 non-null	category
11	fiSecondaryDesc	271971 non-null	category
12	fiModelSeries	58667 non-null	category
13	fiModelDescriptor	74816 non-null	category
14	ProductSize	196093 non-null	category
15	fiProductClassDesc	412698 non-null	category
16	state	412698 non-null	category
17	ProductGroup	412698 non-null	category
18	ProductGroupDesc	412698 non-null	category
19	Drive_System	107087 non-null	category
20	Enclosure	412364 non-null	category
21	Forks	197715 non-null	category
22	Pad_Type	81096 non-null	category
23	Ride_Control	152728 non-null	category
24	Stick	81096 non-null	category
25	Transmission	188007 non-null	category
26	Turbocharged	81096 non-null	category
27	Blade_Extension	25983 non-null	category
28	Blade_Width	25983 non-null	category
29	Enclosure_Type	25983 non-null	category
30	Engine_Horsepower	25983 non-null	category
31	Hydraulics	330133 non-null	category
32	Pushblock	25983 non-null	category
33	Ripper	106945 non-null	category
34	Scarifier	25994 non-null	category
35	Tip_Control	25983 non-null	category
36	Tire_Size	97638 non-null	category
37	Coupler	220679 non-null	category
38	Coupler_System	44974 non-null	category
39	Grouser_Tracks	44875 non-null	category
40	Hydraulics_Flow	44875 non-null	category
41	Track_Type	102193 non-null	category
42	Undercarriage_Pad_Width	102916 non-null	category
43	Stick_Length	102261 non-null	category
44	Thumb	102332 non-null	category

```

45 Pattern_Changer          102261 non-null category
46 Grouser_Type             102193 non-null category
47 Backhoe_Mounting         80712 non-null category
48 Blade_Type               81875 non-null category
49 Travel_Controls          81877 non-null category
50 Differential_Type        71564 non-null category
51 Steering_Controls        71522 non-null category
52 saleYear                 412698 non-null int64
53 saleMonth                412698 non-null int64
54 saleDay                  412698 non-null int64
55 saleDayOfWeek            412698 non-null int64
56 saleDayOfYear            412698 non-null int64
dtypes: category(44), float64(3), int64(10)
memory usage: 63.2 MB

```

We have successfully changed the datatype to category

```
[29]: pd.api.types.is_string_dtype(df_tmp["state"])
```

```
[29]: False
```

```
[30]: pd.api.types.is_categorical_dtype(df_tmp["state"])
```

```
[30]: True
```

```
[31]: df_tmp["state"].cat.categories
```

```
[31]: Index(['Alabama', 'Alaska', 'Arizona', 'Arkansas', 'California', 'Colorado',
        'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Hawaii', 'Idaho',
        'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana',
        'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota',
        'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada',
        'New Hampshire', 'New Jersey', 'New Mexico', 'New York',
        'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon',
        'Pennsylvania', 'Puerto Rico', 'Rhode Island', 'South Carolina',
        'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah', 'Vermont',
        'Virginia', 'Washington', 'Washington DC', 'West Virginia', 'Wisconsin',
        'Wyoming'],
        dtype='object')
```

```
[32]: df_tmp["state"].cat.codes[:10]
```

```
[32]: 205615    43
      274835     8
      141296     8
      212552     8
      62755     8
      54653     8
```

```

81383      8
204924     8
135376     8
113390     8
dtype: int8

```

```
[33]: df_tmp["state"][:10]
```

```

[33]: 205615      Texas
      274835      Florida
      141296      Florida
      212552      Florida
      62755       Florida
      54653       Florida
      81383       Florida
      204924      Florida
      135376      Florida
      113390      Florida
      Name: state, dtype: category
      Categories (53, object): ['Alabama' < 'Alaska' < 'Arizona' < 'Arkansas' ...
      'Washington DC' < 'West Virginia' < 'Wisconsin' < 'Wyoming']

```

We can see that Texas has a code of 43 while Florida has a code of 8.
The code is set alphabetically.

1.4.5 Saving Preprocessed data

```
[34]: df_tmp.to_csv("data/bluebook-for-bulldozers/train_valid_tmp.csv", index = False)
```

```

[35]: df_tmp = pd.read_csv("data/bluebook-for-bulldozers/train_valid_tmp.csv",
      ↪ low_memory = False)

```

1.4.6 Dealing with missing data

Our dataset contains a lot of missing data

```
[36]: df_tmp.isnull().sum()
```

```

[36]: SalesID              0
      SalePrice            0
      MachineID           0
      ModelID             0
      datasource           0
      auctioneerID        20136
      YearMade             0
      MachineHoursCurrentMeter  265194
      UsageBand           339028
      fiModelDesc          0

```

fiBaseModel	0
fiSecondaryDesc	140727
fiModelSeries	354031
fiModelDescriptor	337882
ProductSize	216605
fiProductClassDesc	0
state	0
ProductGroup	0
ProductGroupDesc	0
Drive_System	305611
Enclosure	334
Forks	214983
Pad_Type	331602
Ride_Control	259970
Stick	331602
Transmission	224691
Turbocharged	331602
Blade_Extension	386715
Blade_Width	386715
Enclosure_Type	386715
Engine_Horsepower	386715
Hydraulics	82565
Pushblock	386715
Ripper	305753
Scarifier	386704
Tip_Control	386715
Tire_Size	315060
Coupler	192019
Coupler_System	367724
Grouser_Tracks	367823
Hydraulics_Flow	367823
Track_Type	310505
Undercarriage_Pad_Width	309782
Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176
saleYear	0
saleMonth	0
saleDay	0
saleDayOfWeek	0
saleDayOfYear	0

dtype: int64

```
[37]: df_tmp.isnull().sum() / len(df_tmp)
```

```
[37]: SalesID          0.000000
      SalePrice       0.000000
      MachineID       0.000000
      ModelID         0.000000
      datasource      0.000000
      auctioneerID    0.048791
      YearMade        0.000000
      MachineHoursCurrentMeter 0.642586
      UsageBand       0.821492
      fiModelDesc     0.000000
      fiBaseModel     0.000000
      fiSecondaryDesc  0.340993
      fiModelSeries   0.857845
      fiModelDescriptor 0.818715
      ProductSize     0.524851
      fiProductClassDesc 0.000000
      state           0.000000
      ProductGroup    0.000000
      ProductGroupDesc 0.000000
      Drive_System    0.740520
      Enclosure       0.000809
      Forks           0.520921
      Pad_Type        0.803498
      Ride_Control    0.629928
      Stick           0.803498
      Transmission    0.544444
      Turbocharged    0.803498
      Blade_Extension 0.937041
      Blade_Width     0.937041
      Enclosure_Type  0.937041
      Engine_Horsepower 0.937041
      Hydraulics      0.200062
      Pushblock       0.937041
      Ripper          0.740864
      Scarifier       0.937014
      Tip_Control     0.937041
      Tire_Size       0.763415
      Coupler         0.465277
      Coupler_System  0.891024
      Grouser_Tracks  0.891264
      Hydraulics_Flow 0.891264
      Track_Type      0.752378
      Undercarriage_Pad_Width 0.750626
```

Stick_Length	0.752213
Thumb	0.752041
Pattern_Changer	0.752213
Grouser_Type	0.752378
Backhoe_Mounting	0.804428
Blade_Type	0.801610
Travel_Controls	0.801606
Differential_Type	0.826595
Steering_Controls	0.826697
saleYear	0.000000
saleMonth	0.000000
saleDay	0.000000
saleDayOfWeek	0.000000
saleDayOfYear	0.000000
dtype:	float64

1.4.7 Filling Features with Numerical Data

Let's first see all the features with numerical data

```
[38]: for col_name, col_content in df_tmp.items():
        if pd.api.types.is_numeric_dtype(col_content):
            print(col_name)
```

```
SalesID
SalePrice
MachineID
ModelID
datasource
auctioneerID
YearMade
MachineHoursCurrentMeter
saleYear
saleMonth
saleDay
saleDayOfWeek
saleDayOfYear
```

Let's check the missing values in numerical features

```
[39]: for col_name, col_content in df_tmp.items():
        if pd.api.types.is_numeric_dtype(col_content):
            if col_content.isnull().sum():
                print(col_name, col_content.isnull().sum())
```

```
auctioneerID 20136
MachineHoursCurrentMeter 265194
```

Let's fill these missing numeric values with median of the feature

```
[44]: for col_name, col_content in df_tmp.items():
        if pd.api.types.is_numeric_dtype(col_content):
            if pd.isnull(col_content).sum():
                df_tmp[col_name] = col_content.fillna(col_content.median())
                df_tmp[col_name + "is_missing"] = pd.isnull(col_content)
                # making new cols to check if data was missing.
```

Let's check for missing values again

```
[45]: for col_name, col_content in df_tmp.items():
        if pd.api.types.is_numeric_dtype(col_content):
            if col_content.isnull().sum():
                print(col_name, col_content.isnull().sum())
```

No missing values left.

Now let's check how many values were filled.

```
[46]: df_tmp["auctioneerIDis_missing"].value_counts()
```

```
[46]: False    392562
      True     20136
      Name: auctioneerIDis_missing, dtype: int64
```

```
[47]: df_tmp["MachineHoursCurrentMeteris_missing"].value_counts()
```

```
[47]: True      265194
      False    147504
      Name: MachineHoursCurrentMeteris_missing, dtype: int64
```

1.4.8 Filling missing categories and turning categorical data into numeric data

```
[53]: pd.Categorical(df_tmp["state"]).codes
```

```
[53]: array([43,  8,  8, ...,  4,  4,  4], dtype=int8)
```

By default missing values in pandas has a category of -1

```
[55]: pd.Categorical(df_tmp["UsageBand"]).codes
```

```
[55]: array([-1, -1, -1, ..., -1, -1, -1], dtype=int8)
```

```
[56]: for col_name, col_content in df_tmp.items():
        if not pd.api.types.is_numeric_dtype(col_content):
            # Creating a binary feature showing if the value was missing
            df_tmp[col_name + "_is_missing"] = col_content.isnull()
            # Turning Categories into Numbers and Filling Null Vals
            df_tmp[col_name] = pd.Categorical(col_content).codes + 1
```

Adding +1 to the category codes will make Null's value from -1 to 0.


```
[57]: df_tmp.isnull().sum()
```

```
[57]: SalesID          0
      SalePrice       0
      MachineID       0
      ModelID         0
      datasource      0
      ..
      Backhoe_Mounting_is_missing  0
      Blade_Type_is_missing       0
      Travel_Controls_is_missing  0
      Differential_Type_is_missing 0
      Steering_Controls_is_missing 0
      Length: 103, dtype: int64
```

All the null values have been filled

```
[58]: df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Columns: 103 entries, SalesID to Steering_Controls_is_missing
dtypes: bool(46), float64(3), int16(4), int64(10), int8(40)
memory usage: 77.9 MB
```

1.5 5. Modelling

Now that we have turned dtypes to Numerical and Filled the missing data, we can start modelling

```
[61]: %%time

np.random.seed(0)

X = df_tmp.drop("SalePrice", axis = 1)
y = df_tmp["SalePrice"]

model = RandomForestRegressor(n_jobs = -1, random_state = 0)

model.fit(X, y)
```

```
CPU times: user 11min 31s, sys: 6.46 s, total: 11min 37s
Wall time: 1min 46s
```

```
[61]: RandomForestRegressor(n_jobs=-1, random_state=0)
```

Let's initially score the model on the same dataset

```
[62]: model.score(X, y)
```

```
[62]: 0.9875764803061743
```

This scoring metric is on the same set that the model was trained on. So it isn't an accurate representation of the model's generalization

1.5.1 Splitting the data into Training and Validation Sets

The validation set is the data for the year 2012

```
[67]: df_train = df_tmp[df_tmp["saleYear"] != 2012]
      df_val = df_tmp[df_tmp["saleYear"] == 2012]
```

```
[68]: len(df_train), len(df_val)
```

```
[68]: (401125, 11573)
```

```
[78]: X_train, y_train = df_train.drop("SalePrice", axis = 1), df_train["SalePrice"]
      X_val, y_val = df_val.drop("SalePrice", axis = 1), df_val["SalePrice"]
```

```
[79]: len(X_train), len(y_train), len(X_val), len(y_val)
```

```
[79]: (401125, 401125, 11573, 11573)
```

1.5.2 Building an evaluation function

```
[83]: from sklearn.metrics import mean_squared_log_error, mean_absolute_error, r2_score

def rmsle(y_true, y_preds):
    """
    Function that returns the Root Mean Squared Log Error (RMSLE) of y_true and y_preds.
    """
    return np.sqrt(mean_squared_log_error(y_true, y_preds))

def eval_model(model, X_train, X_test, y_train, y_test):
    """
    Makes prediction and evaluations on given model based on X_train, X_test, y_train and y_test data.
    """
    train_preds = model.predict(X_train)
    test_preds = model.predict(X_test)

    eval_scores = {
        "Train Mean Absolute Error" : mean_absolute_error(y_train, train_preds),
        "Test Mean Absolute Error" : mean_absolute_error(y_test, test_preds),
        "Train Root Mean Squared Log Error" : rmsle(y_train, train_preds),
        "Test Root Mean Squared Log Error" : rmsle(y_test, test_preds),
        "Train R^2" : r2_score(y_train, train_preds),
        "Test R^2" : r2_score(y_test, test_preds)
```

```
}
return eval_scores
```

1.5.3 Testing our model on a subset of the data (to tune the hyperparameters)

```
[90]: %%time
model = RandomForestRegressor(n_jobs = -1, random_state = 0, max_samples = 10000)
model.fit(X_train, y_train)
```

CPU times: user 22.7 s, sys: 264 ms, total: 22.9 s
Wall time: 3.19 s

```
[90]: RandomForestRegressor(max_samples=10000, n_jobs=-1, random_state=0)
```

```
[91]: baseline_scores = eval_model(model, X_train, X_val, y_train, y_val)
baseline_scores
```

```
[91]: {'Train Mean Absolute Error': 5558.52439820505,
      'Test Mean Absolute Error': 7171.1105391860365,
      'Train Root Mean Squared Log Error': 0.25777132630598937,
      'Test Root Mean Squared Log Error': 0.2925990620389206,
      'Train R^2': 0.8606818966052752,
      'Test R^2': 0.832049867497664}
```

1.5.4 Tuning the hyperparameters using RandomizedSearchCV

```
[96]: %%time

from sklearn.model_selection import RandomizedSearchCV

rs_grid = {
    "n_estimators": np.arange(10,200,20),
    "max_depth": [None, 3, 5, 10],
    "min_samples_split": np.arange(2,20,2),
    "min_samples_leaf": np.arange(1,20,2),
    "max_features": [0.5, 1, "sqrt"],
    "max_samples": [10000]
}

rs_model = RandomizedSearchCV(RandomForestRegressor(n_jobs = -1,
    random_state=0),
                               param_distributions = rs_grid,
                               cv = 5,
                               n_iter = 100,
                               verbose = True)
rs_model.fit(X_train, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits
CPU times: user 5min 18s, sys: 49.6 s, total: 6min 7s
Wall time: 9min 36s

```
[96]: RandomizedSearchCV(cv=5,
                        estimator=RandomForestRegressor(n_jobs=-1, random_state=0),
                        n_iter=100,
                        param_distributions={'max_depth': [None, 3, 5, 10],
                                           'max_features': [0.5, 1, 'sqrt'],
                                           'max_samples': [10000],
                                           'min_samples_leaf': array([ 1,  3,  5,
7,  9, 11, 13, 15, 17, 19]),
                                           'min_samples_split': array([ 2,  4,  6,
8, 10, 12, 14, 16, 18]),
                                           'n_estimators': array([ 10,  30,  50,
70,  90, 110, 130, 150, 170, 190])},
                        verbose=True)
```

```
[97]: rs_model.best_params_
```

```
[97]: {'n_estimators': 130,
      'min_samples_split': 10,
      'min_samples_leaf': 3,
      'max_samples': 10000,
      'max_features': 0.5,
      'max_depth': None}
```

```
[98]: baseline_scores
```

```
[98]: {'Train Mean Absolute Error': 5558.52439820505,
      'Test Mean Absolute Error': 7171.1105391860365,
      'Train Root Mean Squared Log Error': 0.25777132630598937,
      'Test Root Mean Squared Log Error': 0.2925990620389206,
      'Train R^2': 0.8606818966052752,
      'Test R^2': 0.832049867497664}
```

```
[99]: rs_scores = eval_model(rs_model, X_train, X_val, y_train, y_val)
      rs_scores
```

```
[99]: {'Train Mean Absolute Error': 5911.904418786114,
      'Test Mean Absolute Error': 7268.799278416505,
      'Train Root Mean Squared Log Error': 0.27010028208693937,
      'Test Root Mean Squared Log Error': 0.29567630218568863,
      'Train R^2': 0.8437098487721018,
      'Test R^2': 0.8271379483482113}
```

1.5.5 Training a model with the tuned hyperparameters

```
[102]: tuned_model = RandomForestRegressor(n_estimators = 130,
                                         min_samples_split = 10,
                                         min_samples_leaf = 3,
                                         max_samples = None,
                                         max_features = 0.5,
                                         max_depth = None,
                                         n_jobs = -1,
                                         random_state = 0)

tuned_model.fit(X_train, y_train)
```

```
[102]: RandomForestRegressor(max_features=0.5, min_samples_leaf=3,
                             min_samples_split=10, n_estimators=130, n_jobs=-1,
                             random_state=0)
```

```
[104]: tuned_model_scores = eval_model(tuned_model, X_train, X_val, y_train, y_val)
tuned_model_scores
```

```
[104]: {'Train Mean Absolute Error': 2844.8336474314638,
        'Test Mean Absolute Error': 5884.576654930028,
        'Train Root Mean Squared Log Error': 0.1423311430882674,
        'Test Root Mean Squared Log Error': 0.2417827606636782,
        'Train R^2': 0.9597602190360334,
        'Test R^2': 0.8834509262663633}
```

1.5.6 Making Prediction on the Test Data

```
[122]: df_test = pd.read_csv("data/bluebook-for-bulldozers/Test.csv",
                             low_memory=False, parse_dates=["saledate"])
df_test.head()
```

```
[122]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	\
0	1227829	1006309	3168	121	3	1999	
1	1227844	1022817	7271	121	3	1000	
2	1227847	1031560	22805	121	3	2004	
3	1227848	56204	1269	121	3	2006	
4	1227863	1053887	22312	121	3	2005	

	MachineHoursCurrentMeter	UsageBand	saledate	fiModelDesc	...	\
0	3688.0	Low	2012-05-03	580G	...	
1	28555.0	High	2012-05-10	936	...	
2	6038.0	Medium	2012-05-10	EC210BLC	...	
3	8940.0	High	2012-05-10	330CL	...	
4	2286.0	Low	2012-05-10	650K	...	

	Undercarriage_Pad_Width	Stick_Length	Thumb	Pattern_Changer	\
0	NaN	NaN	NaN	NaN	

1		NaN		NaN		NaN
2	None or Unspecified		9' 6"	Manual	None or Unspecified	
3	None or Unspecified	None or Unspecified		Manual		Yes
4		NaN		NaN		NaN

	Grouser_Type	Backhoe_Mounting	Blade_Type	Travel_Controls	\
0	NaN	NaN	NaN		NaN
1	NaN	NaN	NaN		NaN
2	Double	NaN	NaN		NaN
3	Triple	NaN	NaN		NaN
4	NaN	None or Unspecified	PAT	None or Unspecified	

	Differential_Type	Steering_Controls
0	NaN	NaN
1	Standard	Conventional
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

[5 rows x 52 columns]

The test data must first be preprocessed to be in the same format as that of our training set

```
[123]: X_train.columns
```

```
[123]: Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
            'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'fiModelDesc',
            'fiBaseModel',
            ...,
            'Undercarriage_Pad_Width_is_missing', 'Stick_Length_is_missing',
            'Thumb_is_missing', 'Pattern_Changer_is_missing',
            'Grouser_Type_is_missing', 'Backhoe_Mounting_is_missing',
            'Blade_Type_is_missing', 'Travel_Controls_is_missing',
            'Differential_Type_is_missing', 'Steering_Controls_is_missing'],
            dtype='object', length=102)
```

```
[124]: df_test.columns
```

```
[124]: Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
            'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'saledate',
            'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc', 'fiModelSeries',
            'fiModelDescriptor', 'ProductSize', 'fiProductClassDesc', 'state',
            'ProductGroup', 'ProductGroupDesc', 'Drive_System', 'Enclosure',
            'Forks', 'Pad_Type', 'Ride_Control', 'Stick', 'Transmission',
            'Turbocharged', 'Blade_Extension', 'Blade_Width', 'Enclosure_Type',
            'Engine_Horsepower', 'Hydraulics', 'Pushblock', 'Ripper', 'Scarifier',
            'Tip_Control', 'Tire_Size', 'Coupler', 'Coupler_System',
            'Grouser_Tracks', 'Hydraulics_Flow', 'Track_Type',
```

```

'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb', 'Pattern_Changer',
'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type', 'Travel_Controls',
'Differential_Type', 'Steering_Controls'],
dtype='object')

```

1.5.7 Preprocessing the data

```

[125]: def preprocess_data(df):

    # Feature Engineering using saledate feature
    df["saleYear"] = df["saledate"].dt.year
    df["saleMonth"] = df["saledate"].dt.month
    df["saleDay"] = df["saledate"].dt.day
    df["saleDayOfWeek"] = df["saledate"].dt.dayofweek
    df["saleDayOfYear"] = df["saledate"].dt.dayofyear
    df.drop("saledate", axis = 1, inplace = True)

    # Converting String data to Categories
    for col_name, col_content in df.items():
        if pd.api.types.is_string_dtype(col_content):
            df[col_name] = col_content.astype("category").cat.as_ordered()

    # Filling data
    for col_name, col_content in df.items():

        # Numerical Data
        if pd.api.types.is_numeric_dtype(col_content):
            if pd.isnull(col_content).sum():
                df[col_name] = col_content.fillna(col_content.median())
                df[col_name + "_is_missing"] = pd.isnull(col_content)
                # making new cols to check if data was missing.

        # Categorical Data
        if not pd.api.types.is_numeric_dtype(col_content):
            # Creating a binary feature showing if the value was missing
            df[col_name + "_is_missing"] = col_content.isnull()
            # Turning Categories into Numbers and Filling Null Vals
            df[col_name] = pd.Categorical(col_content).codes + 1

    return df

```

```

[126]: df_test = preprocess_data(df_test)
df_test.head()

```

```

[126]:
SalesID  MachineID  ModelID  datasource  auctioneerID  YearMade  \
0  1227829      1006309      3168          121           3        1999
1  1227844      1022817      7271          121           3        1000

```

2	1227847	1031560	22805	121	3	2004
3	1227848	56204	1269	121	3	2006
4	1227863	1053887	22312	121	3	2005

	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBaseModel	...	\
0	3688.0	2	499	180	...	
1	28555.0	1	831	292	...	
2	6038.0	3	1177	404	...	
3	8940.0	1	287	113	...	
4	2286.0	2	566	196	...	

	Undercarriage_Pad_Width_is_missing	Stick_Length_is_missing	\
0	True	True	
1	True	True	
2	False	False	
3	False	False	
4	True	True	

	Thumb_is_missing	Pattern_Changer_is_missing	Grouser_Type_is_missing	\
0	True	True	True	
1	True	True	True	
2	False	False	False	
3	False	False	False	
4	True	True	True	

	Backhoe_Mounting_is_missing	Blade_Type_is_missing	\
0	True	True	
1	True	True	
2	True	True	
3	True	True	
4	False	False	

	Travel_Controls_is_missing	Differential_Type_is_missing	\
0	True	True	
1	True	False	
2	True	True	
3	True	True	
4	False	True	

	Steering_Controls_is_missing
0	True
1	False
2	True
3	True
4	True

[5 rows x 101 columns]


```
[127]: len(df_test.columns), len(X_train.columns)
```

```
[127]: (101, 102)
```

```
[129]: set(X_train.columns) - set(df_test.columns)
```

```
[129]: {'auctioneerIDis_missing'}
```

There is no auctioneerID Null values in Test Dataset hence this feature is missing. We can manually add the feature

```
[130]: df_test["auctioneerIDis_missing"] = False
```

```
[139]: df_test = df_test.reindex(columns=list(X_train.columns))
```

```
[140]: len(df_test.columns), len(X_train.columns)
```

```
[140]: (102, 102)
```

1.5.8 Making Predictions on the Test set

```
[141]: test_preds = tuned_model.predict(df_test)
```

```
[142]: test_preds
```

```
[142]: array([17347.62069212, 18622.20832936, 49162.08708438, ...,  
        12535.01052965, 17110.6827157 , 28608.18839628])
```

```
[143]: test_preds_df = pd.DataFrame()  
test_preds_df["SalesID"] = df_test["SalesID"]  
test_preds_df["SalesPrice"] = test_preds
```

```
[144]: test_preds_df
```

```
[144]:
```

	SalesID	SalesPrice
0	1227829	17347.620692
1	1227844	18622.208329
2	1227847	49162.087084
3	1227848	68131.719246
4	1227863	54513.445507
...
12452	6643171	39808.592619
12453	6643173	12412.197032
12454	6643184	12535.010530
12455	6643186	17110.682716
12456	6643196	28608.188396

```
[12457 rows x 2 columns]
```

```
[145]: test_preds_df.to_csv("data/bluebook-for-bulldozers/predicted_sales_price.csv")
```

1.5.9 Feature Importance

```
[146]: tuned_model.feature_importances_
```

```
[146]: array([3.53002628e-02, 1.87088656e-02, 4.52479013e-02, 1.69110698e-03,
        3.31467160e-03, 2.02627988e-01, 3.03853416e-03, 1.06132675e-03,
        4.36857562e-02, 4.82178584e-02, 6.51843745e-02, 4.59642257e-03,
        1.65203621e-02, 1.54674399e-01, 4.05789591e-02, 6.16942667e-03,
        4.80438342e-03, 1.99096712e-03, 3.24624206e-03, 6.15668687e-02,
        5.12113586e-04, 1.57350151e-04, 8.97199123e-04, 1.72662045e-04,
        1.28939126e-03, 1.73156754e-05, 1.57434491e-03, 8.80231196e-03,
        2.74317863e-03, 1.18944838e-03, 4.88483269e-03, 2.81428978e-03,
        3.32742596e-03, 1.04646077e-03, 1.54317753e-03, 7.08262326e-03,
        8.90458833e-04, 1.01528916e-02, 1.97574037e-03, 2.80213202e-03,
        1.11086498e-03, 1.00593958e-03, 2.27993834e-03, 6.42687057e-04,
        6.23172487e-04, 3.74938015e-04, 4.75575766e-04, 2.21729496e-03,
        8.54676772e-04, 2.84932304e-04, 2.18368761e-04, 7.27219199e-02,
        4.25316178e-03, 6.1113228e-03, 3.09441578e-03, 1.01481018e-02,
        1.94622167e-04, 1.39379163e-03, 3.91264434e-04, 0.00000000e+00,
        0.00000000e+00, 2.69055831e-03, 1.32994136e-03, 5.85428315e-03,
        2.98569168e-02, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 6.20083472e-05, 2.01636053e-06, 1.90010416e-04,
        5.17599255e-06, 1.47236631e-04, 4.46252777e-06, 2.46546900e-04,
        2.14456025e-05, 1.28003835e-03, 2.09497500e-03, 1.62409227e-03,
        1.41079037e-03, 1.89688471e-03, 2.68196484e-03, 1.40299783e-03,
        3.40674675e-04, 8.94960103e-04, 3.26197466e-03, 1.62380902e-04,
        1.29479445e-02, 1.84729085e-03, 2.08935085e-03, 4.30793411e-05,
        6.66216316e-05, 6.92148014e-05, 3.88438616e-05, 4.72939387e-05,
        4.84143312e-05, 3.59376961e-04, 1.90987770e-04, 1.01040481e-04,
        8.28277842e-05, 1.06552600e-04])
```

```
[147]: def plot_features(cols, importance, n = 20):
        df = (pd.DataFrame({"features" : cols,
                           "Feature Importance" : importance})).
            sort_values("Feature Importance", "ascending").
            reset_index(drop = True))

        #plotting
        fig, ax = plt.subplots()
        ax.barh(df["features"][:n], df["Feature Importance"][:n])
        ax.set_ylabel("Features")
        ax.set_xlabel("Feature Importance")
        ax.invert_yaxis();
```

```
[151]: plot_features(X_train.columns, tuned_model.feature_importances_)
```

